

Highlights

Structurally-Calibrated Functional Attribution: A Methodology for Process Supervision Weighting in Reflective Reasoning

Haoran Ma, Ningning Wang

- SC-FMA calibrates interventional utility into supervision weights.
- QP is strongest only on the controlled synthetic proxy benchmark.
- PRM800K real evidence is led by `w_struct`; Ridge is closest.
- Audit-prioritization evidence is offline locked-split context only.
- Failed GSM8K/HotpotQA and downstream routes are kept out of claims.

Structurally-Calibrated Functional Attribution: A Methodology for Process Supervision Weighting in Reflective Reasoning

Haoran Ma^a, Ningning Wang^{a,b,*}

^aCollege of Management Science and Engineering, Beijing Information Science and Technology University, Beijing, 102200, China

^bInstitute of Information Systems, ESG Intelligent Application Innovation Research Center, Beijing, 102200, China

ARTICLE INFO

Keywords:

structural calibration
functional attribution
process supervision
convex optimization
reflective cognition
step importance ranking

ABSTRACT

Knowledge-intensive reasoning systems increasingly rely on reflective steps such as self-evaluation, error diagnosis, and plan revision to organize verification workflows. This paper presents Structurally-Calibrated Functional Attribution (SC-FMA), a methodology for auditable verification-step weighting in knowledge-intensive reasoning systems. SC-FMA converts local functional utility signals into calibrated supervision weights through the Structurally-Calibrated Utility (SCU) objective, a convex optimization problem that combines fidelity to utility, graph-level structural necessity, redundancy reduction, and bottleneck protection. On a controlled synthetic benchmark with proxy step-importance labels, the full Quadratic Program (QP) variant is strongest ($\rho = 0.608$) and improves over raw CIU ($\rho = 0.483$). On PRM800K real process-supervision annotations, the primary positive evidence is instead the stored `w_struct` step-ranking signal ($\rho = 0.611$); Ridge is the closest SC-FMA approximation ($\rho = 0.604$), while QP ($\rho = 0.442$) and Projection ($\rho = -0.135$) are downgraded for this real-data route. A new PRM800K audit-prioritization readout shows that the same locked-split evidence can concentrate high-rated process steps under a fixed review budget, but only as offline audit context. GSM8K/HotpotQA task-specific replay, downstream filtering, PRM training, and out-of-distribution transfer remain unsupported. The contribution is therefore a claim-bounded calibration methodology for KBS-oriented verification-step weighting, not a deployed training or replay validation result.

1. Introduction

Knowledge-intensive reasoning systems require verification workflows whose intermediate checks are both locally useful and structurally well placed. In expert systems, knowledge graphs, and reflective language-agent pipelines, a verification step may correlate with final-answer success while still being redundant because another rule, entity path, or neighboring reflection supplies equivalent support. Conversely, a rare bottleneck check may carry only a weak local utility signal but still gate downstream inference. This makes process supervision a Knowledge-Based Systems (KBS) problem as well as a scoring problem: supervision weights should reflect local utility, structural dependency, redundancy, and bottleneck protection.

Existing process-supervision approaches usually treat step correctness, verifier score, or local intervention response as the supervision target [3, 8, 9]. Those signals are useful, but they do not distinguish between a locally helpful step and a structurally necessary one. Prior diagnostic FMA results show that local utility and structural necessity can be weakly aligned, so direct utility weighting can over-weight fluent but redundant reflection and under-weight structurally critical operations.

We propose **Structurally-Calibrated Functional Attribution (SC-FMA)**. SC-FMA takes local functional attribution signals and produces calibrated supervision weights by solving a constrained optimization problem over a reasoning graph. The method is designed for auditable verification-step weighting: all inputs are explicit, all calibration variants are deterministic, and claim boundaries are tied to stored artifacts.

The paper makes five contributions:

1. a convex Structurally-Calibrated Utility (SCU) objective that combines local utility, graph necessity, redundancy, and bottleneck protection;

*Corresponding author.

 mahaoran0000@foamail.com (H. Ma); wangningning@bistu.edu.cn (N. Wang)

ORCID(s):

2. three deterministic calibration variants—Ridge, Quadratic Program (QP), and Topology Projection—covering different compute budgets;
3. formal guarantees for convexity/uniqueness, non-redundant monotonicity, variance shrinkage, and bottleneck protection;
4. controlled synthetic evidence showing QP as the strongest variant under controlled proxy labels;
5. real PRM800K evidence showing that w_{struct} , not QP, is the main real-data step-ranking result, with Ridge acting as the closest SC-FMA approximation;
6. a bounded PRM800K audit-prioritization readout that treats step weights as review-budget guidance rather than downstream PRM-training evidence.

This evidence boundary is intentional. The synthetic benchmark uses controlled proxy labels, not external ground truth. PRM800K supports step-label ranking and an in-distribution frozen PRM baseline comparison. GSM8K/HotpotQA task-specific replay, downstream filtering, and PRM training are not positive evidence in this submission.

2. Related Work

Process supervision and PRMs. Process Reward Models assign scalar estimates to intermediate reasoning steps, enabling verifier-guided search and process-level feedback [3, 8]. ProcessBench and related evaluations show that step-level supervision remains brittle under distribution shift [9]. SC-FMA is not a PRM-training result. It is a pre-calibration methodology that can transform existing utility or process signals into structure-aware weights before any downstream use.

Attribution for reasoning. Token-level attribution methods such as LIME, Integrated Gradients, and SHAP [4, 6, 7] explain model predictions at fine granularity, but reflective operations are multi-token functional units. SC-FMA works at the step level and calibrates functional attribution signals with graph structure.

Graph reasoning and KBS. Knowledge-graph reasoning emphasizes dependencies among entities, rules, and paths [1]. Classical KBS and cognitive-system work likewise treats inference as structured knowledge manipulation rather than isolated scoring. SC-FMA imports this structural view into reflective reasoning traces by treating verification steps as graph nodes with necessity and redundancy diagnostics.

Convex calibration. Convex optimization provides globally analyzable objectives and stable solvers [2, 5]. SC-FMA uses this tradition to make process-supervision weighting auditable: a calibrated weight is the solution to an explicit objective, not a heuristic post-processing rule.

3. Methodology

3.1. Local Utility and Intervention Boundary

Let a reasoning trace contain k reflective steps $s_1, \dots, s_k$. The local utility vector $\mathbf{c}$ records how much each step contributes under the documented intervention protocol. In the primary CIU route, interventions are length-preserving MASK replays: the targeted span is replaced with mask tokens while preserving token count and position. Trace-level binary correctness and bounded continuous utility/outcome fields in $[0, 1]$ are supported. Fields named `correctness` remain binary by contract.

DELETE, SHUFFLE, TRUNCATE, and CONTRADICT perturbations are retained only as destructive diagnostics and robustness checks. They are useful for stress-testing structural behavior, but they are not structure-preserving primary CIU interventions. The implementation records this boundary in intervention metadata with `structure_preserving=false`, `primary_ciu_allowed=false`, and `proxy_only=true`.

3.2. Reflection Graph

SC-FMA constructs a reflection graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ whose nodes are reflective steps. Edges encode documented structural relations such as temporal adjacency and topical similarity. From this graph, the method computes:

- structural necessity $\mathbf{n}$ from graph perturbation diagnostics;
- redundancy matrix $\mathbf{R}$ from pairwise substitutability indicators;
- bottleneck indicator $\mathbf{b}$ for rare steps whose removal strongly affects downstream structure.

The graph is an operational approximation, not a claim of recovered causal structure.

3.3. SCU Objective

The calibrated weight vector $\mathbf{w}$ lies on the probability simplex. SC-FMA solves:

$$L(\mathbf{w}) = \alpha \|\mathbf{w} - \tilde{\mathbf{c}}\|_2^2 + \beta \|\mathbf{w} - \tilde{\mathbf{n}}\|_2^2 + \gamma \mathbf{w}^\top \mathbf{R} \mathbf{w} - \delta \sum_i b_i \log(w_i), \quad (1)$$

subject to $w_i \geq \epsilon$ and $\sum_i w_i = 1$. Here $\tilde{\mathbf{c}}$ and $\tilde{\mathbf{n}}$ are normalized local utility and structural necessity. The first term preserves fidelity to local utility; the second aligns weights with graph necessity; the third penalizes redundant high weights; and the fourth protects bottlenecks from collapse.

3.4. Calibration Variants

Ridge uses a closed-form softmax over a linear combination of local utility and structural necessity. It is the simplest and fastest route.

QP solves the full SCU objective with redundancy and bottleneck terms. It is the strongest synthetic-benchmark variant but is not the best PRM800K real-data variant.

Projection applies a fast topology-constrained projection. It is included as a computationally cheap ablation and is not supported as the primary PRM800K result.

Conditional replacement, counterfactual matching, and doubly robust estimation are planned extensions rather than completed evidence in this submission.

3.5. Guarantees

Under positive regularization and $w_i > \epsilon$, the SCU objective is strictly convex on the feasible simplex and has a unique global solution. For non-redundant step pairs with identical redundancy profiles, the calibrated weights preserve the joint ordering implied by $\tilde{\mathbf{c}}$ and $\tilde{\mathbf{n}}$. The structural term shrinks weight variance relative to raw utility weighting when structural necessity has lower variance than CIU, and the bottleneck log barrier prevents identified bottleneck nodes from receiving zero weight. Full derivations, KKT conditions, and solver diagnostics are provided in the supplementary material.

4. KBS Fit: Verification-Step Weighting as Structured Decision Support

SC-FMA frames verification-step weighting as a knowledge-structured decision-support problem. A reflective reasoning trace is not only a sequence of locally scored steps: it is also a graph of dependencies, redundancies, and bottleneck checks. This structure is familiar to KBS settings in which rules, entity paths, ontology relations, and verification routines interact rather than contribute independently.

The methodology therefore exposes three KBS-relevant objects. First, graph dependencies make the weighting decision inspectable at the step and edge level. Second, redundancy diagnostics identify when several high-scoring steps provide overlapping support. Third, bottleneck protection prevents rare but structurally central checks from being washed out by smooth utility scores. These objects make the output suitable for audit and review-budget allocation, even before any downstream training loop is attempted.

The ontology-aware edge pilot remains deliberately narrow. It records ontology relation labels as fixture-level diagnostic metadata and shows how ontology-derived edges could replace purely topical edges in a future domain-specific KBS workflow. It does not validate a rule engine, ontology reasoner, KG query engine, or deployed knowledge-base workflow. Such validation would require a real knowledge-base task, an executable reasoning component, and a preregistered evaluation protocol.

5. Experiments

5.1. Data and Claim Permissions

Table 1 summarizes the evidence routes used in this submission. The distinction between controlled synthetic evidence and PRM800K real-data evidence is central to the claim boundary.

5.2. Baseline and Fairness Ledger

The main comparison claim is deliberately narrow: SC-FMA-style weighting should be judged against same-split step-ranking controls and protocol-compatible variants, not against unrelated PRM training pipelines. Table 2 records the fairness status of each comparator used or discussed in the submission.

Table 1

Evidence surfaces, claim permissions, and forbidden upgrades.

Evidence surface	Status	Allowed claim	Forbidden upgrade
Synthetic controlled benchmark	Passed	Method calibration and controlled proxy-label ranking; QP is strongest only here.	External step-importance truth or real-task superiority.
PRM800K v3.6 locked step-label ranking	Passed	Real PRM800K step-label ranking with w_struct as the primary signal.	Downstream PRM training, filtering, or task replay validation.
PRM800K v3.8 frozen PRM baseline	Passed, context-only	In-distribution frozen PRM baseline context on the same PRM800K locked split.	Clean non-overlap public-PRM comparison or cross-dataset generalization.
PRM800K audit prioritization	Added, context-only	Offline review-budget readout over the same locked PRM800K split.	Operational deployment, PRM training gain, or filtering superiority.
Stage 2 holdout	Diagnostic	Small, stratum-dependent diagnostic signal only ($\rho = 0.1628$).	Broad real-task validation.
GSM8K/HotpotQA replay and filtering	Failed or blocked	Negative provenance and route boundary.	Positive replay, downstream filtering, or PRM-training evidence.
Ontology-aware edge pilot	Fixture-level diagnostic	KBS extension path for ontology-derived graph metadata.	Rule-engine, ontology-reasoner, KG-query, or deployed KBS validation.

Table 2

Baseline and fairness ledger for reviewer-facing comparison boundaries.

Comparator	Role	Data/split	Fairness status	Main use
raw_local_utility	Direct ablation	Same split PRM800K locked	Pass: same rows, labels, and metric	Shows raw utility alone is insufficient.
span_length	Trivial control	Same split PRM800K locked	Pass: deterministic same-split control	Tests length-driven explanation.
relative_position	Trivial control	Same split PRM800K locked	Pass: deterministic same-split control	Tests position-driven explanation.
random	Sanity control	Same split PRM800K locked	Pass: deterministic seeded control	Calibrates chance ranking.
Frozen PRM prefix score	In-distribution context	Same split PRM800K locked	Overlap-limited context only	Provides PRM-distribution reference, not a clean SOTA comparison.
SC-FMA Ridge	Recommended approximation	Same split PRM800K locked	Pass: same features and scoring protocol	Closest SC-FMA approximation to w_struct .
SC-FMA QP	Controlled-synthetic strongest variant	Same split PRM800K locked for downgrade	Pass as variant; downgraded on PRM800K	Shows synthetic strength does not transfer to this real-data route.
Math-Shepherd/public literature	Citation-only PRM	Not protocol-compatible here	Citation-only unless clean non-overlap scoring is added	Positions against PRM literature without unfair numeric claims.
ProcessBench	Citation-only benchmark context	Separate benchmark	Citation-only unless a future route is preregistered	Motivates process-supervision brittleness and future evaluation.

Table 3

Synthetic controlled step-ranking results. Mean metrics across 200 traces.

Method	Spearman ρ	Kendall τ	NDCG@3	Top-3 overlap
SC-FMA QP	0.608	0.533	0.921	0.958
SC-FMA Ridge	0.536	0.461	0.884	0.912
Monte Carlo Shapley	0.524	0.447	0.871	0.901
Gradient×Input	0.491	0.419	0.845	0.883
Raw CIU	0.483	0.411	0.839	0.875
Token surprisal	0.310	0.251	0.702	0.741

5.3. Synthetic Controlled Benchmark

The synthetic benchmark contains 200 traces and 1,027 reflective steps generated by deterministic scripts. The target labels are controlled oracle/proxy step-importance labels designed for method comparison; they are not external ground truth. The oracle-independence audit in the supplementary material verifies low overlap between the proxy-label components and SCU inputs.

Table 4

PRM800K locked split: real step-label ranking and variant boundary.

Signal	Spearman ρ	Δ vs. w_struct	Interpretation
w_struct	0.611	—	Primary real-data step-ranking result.
SC-FMA Ridge	0.604	−0.007	Closest SC-FMA approximation; preserves ranking.
SC-FMA QP	0.442	−0.169	Underperforms on PRM800K; not the real-data winner.
SC-FMA Projection	−0.135	−0.746	Fails on this route.
Frozen PRM prefix score	0.252	−0.360	In-distribution baseline context only.
Raw local utility	−0.077	−0.689	Uncalibrated utility is insufficient.

Table 5

PRM800K locked split as offline audit-prioritization context.

Method	Top-1 max-label hit	Label mass@25%	NDCG@25%	Claim permission
w_struct	0.9169	0.3809	0.9506	Audit-prioritization context only.
SC-FMA Ridge	0.9054	0.3796	0.9460	Audit-prioritization context only.
SC-FMA QP	0.9033	0.3790	0.9439	Audit-prioritization context only; not the real-data ranking winner.
Frozen PRM prefix score	0.7775	0.3419	0.8465	In-distribution baseline context only.
Random	0.7329	0.3065	0.7757	Sanity control.
Span length	0.7055	0.2969	0.7526	Trivial control.

On this controlled benchmark, QP is the strongest variant. It improves over raw CIU by +0.125 Spearman and the ablation study shows that structure alignment, redundancy penalty, and bottleneck protection each contribute to ranking quality. These results support the SCU optimization design under controlled conditions.

5.4. PRM800K Real Step-Label Ranking

The real-data route uses PRM800K phase-2 process-supervision annotations under a frozen hash split: 4,417 samples and 34,219 labeled steps. The primary result is w_struct , which reaches Spearman $\rho = 0.611$ and substantially exceeds raw local utility ($\rho = -0.077$). The frozen PRM baseline comparison is reported only as in-distribution context because it uses the same PRM800K distribution.

Two conclusions follow. First, the real-data main result is w_struct , not QP. Second, Ridge is the only SC-FMA variant that closely preserves the stored real-data ordering. The more complex QP and Projection variants are useful for synthetic calibration analysis but must be downgraded for PRM800K.

5.5. Offline Audit-Prioritization Readout

The same locked PRM800K split can also be read operationally as an audit-prioritization setting: if a reviewer or verifier can inspect only a fixed fraction of process steps, do the scores concentrate high-rated annotations? Table 5 reports this readout for the top methods by NDCG@25%. The result remains a step-ranking use case on the locked PRM800K split. It is not downstream PRM training, not filtering superiority, and not GSM8K/HotpotQA replay evidence.

In this bounded readout, w_struct and SC-FMA Ridge exceed $raw_local_utility$ (NDCG@25% = 0.7221), $relative_position$ (0.3204), and the best simple control, $random$ (0.7757). The result supports the operational statement that high-rated PRM800K process steps can be concentrated under a fixed review budget. It does not support claims about trained PRM improvement, task success, or external datasets.

5.6. Stage 2 Diagnostic Holdout

The Stage 2 holdout result is intentionally reported as diagnostic. The aggregate Spearman signal is small ($\rho = 0.1628$) and stratum-dependent; high-divergence strata show stronger alignment than mid or random strata. This supports the interpretation that SC-FMA is most informative where local utility and structural necessity diverge, but it is not sufficient evidence for broad real-task validation.

6. Discussion

The empirical pattern is useful precisely because it is not uniform. QP is strongest when the benchmark is controlled and proxy labels are aligned with the full SCU objective. PRM800K instead favors w_{struct} , with Ridge preserving that ordering more faithfully than QP or Projection. This separation prevents the paper from turning a synthetic calibration win into a real-data claim it cannot support.

For KBS research, the most durable contribution is the audit surface. The method exposes which verification steps receive weight, which graph relations explain structural necessity, and which redundant or bottleneck relations shape the final allocation. The ontology-aware pilot shows an interface path for adding domain knowledge to that graph, but remains fixture-level diagnostic evidence until a rule engine, ontology reasoner, or KG query workflow is evaluated end to end.

7. Scope and Limitations

The current evidence supports SC-FMA as a diagnostic and calibration methodology, not as a complete downstream training system.

- **Non-identifiability.** The intervention notation is operational shorthand for documented trace manipulation. The framework reports protocol-defined functional attribution scores, not formally identified treatment-effect parameters or global causal quantities.
- **Synthetic scope.** Synthetic results use controlled proxy labels. They justify method behavior under controlled conditions, not external truth.
- **Real-data boundary.** PRM800K supports step-label ranking and in-distribution frozen PRM baseline context. GSM8K/HotpotQA replay and filtering routes remain failed or blocked.
- **Audit-prioritization boundary.** The audit-prioritization readout is an offline locked-split analysis of review-budget concentration. It is not evidence of PRM training, filtering gains, or deployed workflow performance.
- **Intervention boundary.** MASK is the primary structure-preserving CIU intervention. DELETE, SHUFFLE, TRUNCATE, and CONTRADICT are destructive diagnostics only.
- **Planned components.** Conditional replacement, counterfactual matching, and doubly robust estimation remain future work and are not counted as completed evidence.
- **KBS deployment.** No production rule engine, ontology reasoner, KG query engine, or deployed knowledge-base workflow is validated here.

8. Conclusion

SC-FMA converts local functional attribution signals into structurally calibrated supervision weights through a convex SCU objective. The controlled synthetic benchmark supports QP as the strongest method under controlled proxy labels. The PRM800K real-data route supports a different boundary: w_{struct} is the main real step-ranking signal, Ridge closely preserves it, and QP/Projection must be downgraded. The offline audit-prioritization readout adds a bounded KBS-facing use case for review-budget allocation on the same locked split. With these boundaries, SC-FMA offers an auditable methodology for verification-step weighting in reflective reasoning and KBS-oriented analysis, while leaving downstream PRM training, task-specific replay validation, and deployed knowledge-base workflows to future work.

Acknowledgments

The authors thank the College of Management Science and Engineering at Beijing Information Science and Technology University for research support.

Funding

This work was supported by the National Social Science Fund of China Project (24BSH018) and the Beijing Natural Science Foundation Project (L252145).

Declaration of Competing Interest

The authors declared that they have no conflicts of interest to this work.

Declaration of generative AI and AI-assisted technologies in the writing process

During preparation of this work, the authors used OpenAI Codex for language editing, LaTeX packaging assistance, and compliance checks. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

Data Availability

PRM800K is publicly available from its original source. Derived locked-split reports, audit-prioritization artifacts, and reproduction scripts will be made available by the authors on request.

CRedit authorship contribution statement

Haoran Ma: Conceptualization, Methodology, Software, Formal analysis, Investigation, Data Curation, Writing – Original Draft, Visualization. **Ningning Wang:** Methodology, Validation, Writing – Review & Editing, Supervision, Project administration, Funding acquisition.

References

- [1] Bellomarini, L., Fayzrakhmanov, R.R., Gottlob, G., Kravchenko, A., Laurenza, E., Nenov, Y., Reissfelder, S., Sallinger, E., Sherkhonov, E., Vahdati, S., Wu, L., 2024. Knowledge graphs and machine learning: A survey. *Data & Knowledge Engineering* 150, 102247. doi:[10.1016/j.datak.2023.102247](https://doi.org/10.1016/j.datak.2023.102247).
- [2] Boyd, S., Vandenberghe, L., 2004. *Convex Optimization*. Cambridge University Press. URL: <https://web.stanford.edu/~boyd/cvxbook/>, doi:[10.1017/CB09780511804441](https://doi.org/10.1017/CB09780511804441).
- [3] Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., Cobbe, K., 2023. Let's verify step by step. *arXiv preprint arXiv:2305.20050* URL: <https://arxiv.org/abs/2305.20050>, doi:[10.48550/arXiv.2305.20050](https://doi.org/10.48550/arXiv.2305.20050).
- [4] Lundberg, S.M., Lee, S.I., 2017. A unified approach to interpreting model predictions, in: *Advances in Neural Information Processing Systems*, pp. 1–11. URL: <https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions>.
- [5] Nocedal, J., Wright, S.J., 2006. *Numerical Optimization*. 2nd ed., Springer. doi:[10.1007/978-0-387-40065-5](https://doi.org/10.1007/978-0-387-40065-5).
- [6] Ribeiro, M.T., Singh, S., Guestrin, C., 2016. "why should i trust you?": Explaining the predictions of any classifier, in: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1135–1144. URL: <https://doi.org/10.1145/2939672.2939778>, doi:[10.1145/2939672.2939778](https://doi.org/10.1145/2939672.2939778).
- [7] Sundararajan, M., Taly, A., Yan, Q., 2017. Axiomatic attribution for deep networks, in: *Proceedings of the 34th International Conference on Machine Learning*, PMLR. pp. 3319–3328. URL: <https://proceedings.mlr.press/v70/sundararajan17a.html>.
- [8] Wang, P., Li, L., Shao, Z., Xu, R., Dai, D., Li, Y., Chen, D., Wu, Y., Sui, Z., 2024. Math-shepherd: Verify and reinforce LLMs step-by-step without human annotations, in: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics, Bangkok, Thailand. pp. 9426–9439. URL: <https://aclanthology.org/2024.acl-long.510/>, doi:[10.18653/v1/2024.acl-long.510](https://doi.org/10.18653/v1/2024.acl-long.510).
- [9] Zheng, C., Zhang, P., Ding, Z., Zhang, Y., Zhuang, F., Zhou, J., 2024. ProcessBench: A benchmark for identifying errors in mathematical reasoning of language models. *arXiv preprint arXiv:2412.06559* URL: <https://arxiv.org/abs/2412.06559>.