

Supplementary Material for Structurally-Calibrated Functional Attribution for Audit Prioritization in Knowledge-Intensive Reasoning

Haoran Ma; Ningning Wang

Abstract. This supplementary material contains extended proofs, additional controlled-benchmark results, implementation details, and reproducibility notes separated from the main manuscript. The synthetic labels discussed here are controlled proxy labels for method comparison, not external ground truth.

Keywords. supplementary material; structural calibration; reproducibility; knowledge graph pilot

Extended Proofs and Derivations This appendix provides supplementary derivations, edge-case analyses, and tightness constructions referenced in the main text. Full proofs of the variance-reduction and bottleneck-protection theorems, the bottleneck tightness construction, and SLSQP convergence diagnostics are provided in the subsections below.

KKT Conditions for the SCU Objective The Karush-Kuhn-Tucker (KKT) conditions for the SCU constrained optimization problem (Eqs. 1-2) are: [Formula or algorithm omitted in DOCX view; see LaTeX/PDF source for authoritative mathematical layout.] The partial derivative is $L / w_i = 2(w_i - c_i) + 2(w_i - n_i) + 2(Rw)_i - b_i / w_i$. The multiplier $\lambda^* R$ (unrestricted) captures the equality constraint; $\lambda^* \geq 0$ enforce the floor constraints. By complementary slackness, $\lambda^* > 0$ only when $w_i \lambda^* = 0$. The complete derivation and case-by-case analysis of boundary solutions are provided in Section A.1 below.

Derivation of the Variance Reduction Bound The full algebraic derivation is provided below. The key steps are: (i) the two-term objective $L_2(w) = \|w - c\|_2^2 + \|w - n\|_2^2$ has unconstrained minimizer $w = (c + n) / (+)$; (ii) variance decomposes as $\text{Var}(w) = (+)^{-2} \text{Var}(c) + (+)^{-2} \text{Var}(n) + 2 / (+)^2 \text{Cov}(c, n)$; (iii) since $\text{Cov}(c, n) \in [0.05, 0.09]$ and $\text{Var}(n) \leq \text{Var}(c)$, the dominant shrinkage factor is $1 / (+)$; (iv) projection onto the simplex W is non-expansive in variance (wang2013projection), yielding the stated bound.

Tightness Construction for Bottleneck Lower Bound The explicit tightness construction for the bottleneck lower-bound theorem is provided in Section A.3 below. Substituting $w_2 = 1 - w_1$: [Formula or algorithm omitted in DOCX view; see LaTeX/PDF source for authoritative mathematical layout.] The first-order condition: [Formula or algorithm omitted in DOCX view; see LaTeX/PDF source for authoritative mathematical layout.] Setting $\delta = \alpha / (\alpha + \beta)$ makes the floor constraint binding and the lower bound in the bottleneck lower-bound theorem exact.

Monotonicity Violation under Redundancy The monotonicity theorem requires non-redundant steps ($R_{ik} = R_{jk}$ for all k). When redundancy profiles differ, the term deliberately violates monotonicity. We illustrate with a minimal example. Consider $k = 2$ with $c = (0.6, 0.4)^T$, $n = (0.5, 0.5)^T$, $b_1 = b_2 = 0$, and: [Formula or algorithm omitted in DOCX view; see LaTeX/PDF source for authoritative mathematical layout.] Here $c_1 > c_2$ and $n_1 = n_2$, so by CIU alone w_1^* should exceed w_2^* . With $\beta = 1.0$, $\alpha = 0.5$, $\delta = 0.2$, $\gamma = 0$: [Formula or algorithm omitted in DOCX view; see LaTeX/PDF source for authoritative mathematical layout.] Solving numerically yields $w_1^* \approx 0.47$, $w_2^* \approx 0.53$, i.e., $w_2^* > w_1^*$ despite $c_1 > c_2$. The redundancy penalty equalizes the weights because steps 1 and 2 are highly similar ($R_{12} = 0.9$), and this equalization outweighs the fidelity term. This is the designed behavior: structural calibration deliberately overrides local utility ordering when steps are substitutable, preventing the dilution of supervision across redundant operations. Quantitative analysis across the full benchmark: among step pairs with $R_{ij} > 0.7$, monotonicity violations occur in 31.2% of pairs, consistent with the deliberate equalization mechanism.

Additional Experimental Results and Figures This appendix provides extended experimental results, supplementary figures, and detailed stratum analysis that support the main-text findings.

Full Ranking Metrics per Trace Size Table [tab:per-size](#) breaks down the controlled synthetic ranking metrics by trace size ($k = 3$ to 8) for the three SC-FMA variants and raw CIU. Within this controlled proxy-label benchmark, SC-FMA QP is strongest at all k , with the largest gains at $k = 6$ where redundancy effects are most pronounced.

Step importance ranking metrics stratified by trace size k . Mean Spearman over traces of each size. Standard errors in parentheses (bootstrap, 10k resamples). Method | $k=3$ | $k=4$ | $k=5$ | $k=6$ | $k=7$ SC-FMA QP | 0.582 (0.041) | 0.601 (0.038) | 0.615 (0.035) | 0.622 (0.033) | 0.634 (0.029) SC-FMA Ridge | 0.521 (0.044) | 0.534 (0.040) | 0.540 (0.037) | 0.538 (0.035) | 0.542 (0.032) SC-FMA Proj. | 0.328 (0.048) | 0.341 (0.044) | 0.345 (0.041) | 0.339 (0.039) | 0.342 (0.036) Raw CIU | 0.471 (0.045) | 0.485 (0.041) | 0.490 (0.038) | 0.482 (0.036) | 0.486 (0.033) (QP - CIU) | +0.111 | +0.116 | +0.125 | +0.140 | +0.148 [TABLE]

The improvement (QP - CIU) increases monotonically with k (from +0.111 at $k=3$ to +0.148 at $k=7$), consistent with the structural calibration mechanism: larger traces contain more re
[H1] Main-Text Material Moved to Supplementary This appendix records material that was removed from the compressed manuscript to keep the upload manuscript within the project-speci
[H2] Notation and Positioning Table tab:supp-notation provides the compact notation reference used throughout the paper.

SC-FMA is positioned between process supervision and graph-aware audit analysis. Process Reward Models supply step-level supervision signals; attribution methods supply local importance scores; graph diagnostics supply structural constraints. SC-FMA does not train a PRM, certify a KBS deployment, or identify causal effects. It calibrates an existing fidelity signal using observable trace structure.

Structural Necessity and Optimization Algorithms [Formula or algorithm omitted in DOCX view; see [LaTeX/PDF source for authoritative mathematical layout.](#)] [Formula or algorithm omitted in DOCX view; see [LaTeX/PDF source for authoritative mathematical layout.](#)]

Hyperparameters and Ablations Table [tab:supp-hyperparameters](#) lists the hyperparameter search grid and the selected values used for the controlled synthetic benchmark, and [Table tab:supp-ablation](#) reports the corresponding component-contribution diagnostic.

Hyperparameter search grid and selected values used for the controlled synthetic benchmark. Parameter | Search range | Selected value | Role | $\backslash 0.1, 0.5, 1.0, 2.0 \backslash$ | 1.0 | Fidelity to supplied utility anchor. | $\backslash 0.0, 0.25, 0.5, 1.0 \backslash$ | 0.5 | Structural necessity alignment. | $\backslash 0.0, 0.1, 0.2, 0.5, 0.8 \backslash$ | 0.2 | Redundancy penalty. | $\backslash 0.0, 0.05, 0.1, 0.2, 0.4 \backslash$ | 0.1 | Bottleneck log-barrier. | $\backslash 10^{-6}, 10^{-5}, 10^{-4}, 10^{-3} \backslash$ | 10^{-4} | Numerical floor for weights. [TABLE]

QP hyperparameter sensitivity on the structural stress-test generator. Values are five-seed means over 200 traces per seed. is relative to the diagnostic default ($=0.2, =0.1$). This grid is a supplementary QP robustness diagnostic only; it is not positive external validation and does not change the PRM800K Ridge recommendation. $=0.0$ | $=0.05$ | $=0.1$ | $=0.2$ | $=0.4$ 0.0 | 0.4879 (-0.1052) | 0.5413 (-0.0517) | 0.5758 (-0.0173) | 0.5918 (-0.0013) | 0.5841 (-0.0089) 0.1 | 0.4993 (-0.0938) | 0.5556 (-0.0375) | 0.5854 (-0.0076) | 0.6047 (+0.0116) | 0.5961 (+0.0030) 0.2 | 0.5105 (-0.0825) | 0.5632 (-0.0298) | 0.5931 (+0.0000) | 0.6130 (+0.0200) | 0.6055 (+0.0124) 0.5 | 0.5367 (-0.0564) | 0.5894 (-0.0037) | 0.6200 (+0.0269) | 0.6381 (+0.0450) | 0.6260 (+0.0329) 0.8 | 0.5512 (-0.0419) | 0.6048 (+0.0117) | 0.6311 (+0.0380) | 0.6477 (+0.0547) | 0.6383 (+0.0453) [TABLE]

The sensitivity grid shows that the structural stress-test behavior is not caused by a single isolated (,) setting. Removing bottleneck pressure ($=0$) consistently lowers Spearman r
[H2] Efficiency and Scaling Table tab:supp-efficiency reports the per-trace computational cost moved from the main manuscript.

The full scaling figure remains available in the source figure bundle. The main manuscript reports only the page-budget-critical conclusion: Ridge is the lightweight real-data approximation, while QP is used for the controlled synthetic benchmark.

PRM800K Variant Details and Audit Readout : The core PRM800K audit-prioritization readout has been moved to the main manuscript, Section 6 (KBS Audit Demonstration). Tables `tab:supp-scfma-variants` and `tab:supp-audit-prioritization` provide extended variant comparisons and additional baselines that supplement the main-text results.

SC-FMA calibration variant comparison on the PRM800K locked split. All variants use `w_struct` Method | Mean | 95% CI | vs. ``w_struct`` ``w_struct`` | 0.611 | [0.605,0.618] | - SC-FMA Ridge | 0.604 | [0.597,0.611] | -0.007 SC-FMA QP | 0.442 | [0.432,0.452] | -0.169 SC-FMA Projection | -0.135 | [-0.147,-0.122] | -0.746 Raw local utility | -0.077 | [-0.092,-0.064] | -0.689 Frozen PRM prefix score | 0.252 | - | -0.359 Random | 0.006 | [-0.008,0.020] | -0.606 [TABLE]

These tables retain the route-specific downgrade: the strongest real-data result is ``w_struct``; Ridge is a close approximation; QP and Projection are not reported as real-data winners.

WebQSP Fixed-Schema Trace-Audit Diagnostic The WebQSP route is included only as a supplementary reasoning-trace audit diagnostic. It does not evaluate KGQA accuracy, semantic parsing, entity-linking quality, or question-answering model performance. The route converts logical-form-derived WebQSP records into deterministic six-step executable traces: entity linking, relation traversal, candidate generation, candidate verification, ambiguity resolution, and answer verification. Rule replay then masks one reasoning step at a time and computes a replay-derived importance target. The knowledge graph is treated as a knowledge source for trace construction; the evaluation unit is the reasoning step in the trace. Development used the train split and final diagnostic reporting used the test split. The train run produced 494 traces and 2,964 steps. The locked test run produced 1,627 traces and 9,762 steps. The data audit passed on both splits: missing-entity, disconnected-subgraph, duplicate-trace, answer-leakage, invalid-trace, and replay-failure counts were all zero. Replay failure rate is reported descriptively and is not used as a hard pass/fail threshold.

WebQSP fixed-schema separability diagnostic on the locked test split. The route is diagnostic only and is not interpreted as KGQA evidence. Method | Spearman | Pairwise accuracy | NDCG@25% Raw rule delta | 1.000 | 1.000 | 1.000 SC-FMA Ridge | 0.878 | 1.000 | 1.000 SC-FMA QP | 0.526 | 0.799 | 0.808 Relative position | 0.293 | 0.667 | 1.000 Random | -0.007 | 0.496 | 0.582 Graph degree | -0.726 | 0.111 | 0.493 [TABLE]

The diagnostic exposes a structural property of the fixed trace schema. Entity linking, relation traversal, and answer verification are indispensable under the replay rules and rec
[H2] MuSiQue KBS-style Knowledge-Audit Details The MuSiQue route is an offline, deterministic KBS-style audit-prioritization experiment. It constructs trace records from MuSiQue de

The support decision is reported for completeness but is not interpreted as positive evidence. Because the audit labels and scoring features share a common step-type source, the task is close to trivial: a hand-coded step-type priority rule (answer synthesis > bridge binding > evidence check > retrieval/question decomposition > distractor rejection) attains NDCG@25% of 1.0000, Top-1 hit of 1.0000, and Mass@25% of 0.3939 across all 1,415 locked traces, identical to the ``w_struct``, SC-FMA Ridge, and SC-FMA QP rows. The @25% of 0.1193 over graph centrality (0.8807), with bootstrap 95% CI [0.1187, 0.1199], therefore reflects the gap between methods that use the step-type priority and controls that do not, not a calibration advantage; the narrow interval is itself a consequence of the label structure (only three distinct label vectors across the 1,415 locked traces). The field-level leakage

audit is clean--locked labels are not read during scoring, forbidden fields are excluded, and label permutation reduces the mean diagnostic delta from 0.1347 to 0.0078--but field-level checks cannot detect construction-level homology. The route is therefore a feasibility demonstration only; it is not production KBS deployment validation and does not provide independent audit-prioritization evidence.

KBS Integration Notes and KG Fixture Pilot : The KBS methodological analogy, including the detailed mapping of high-necessity, high-redundancy, and bottleneck steps to KBS review concepts, is discussed in the main manuscript's KBS methodological analogy subsection. The core KBS implications and limitations are discussed in the main manuscript, Section 5. The fixture-level KG pilot content below extends the main-text discussion with implementation details that are not covered in the main manuscript. The fixture-level KG pilot uses the Countries knowledge graph with country and region entities and ``locatedIn`/`neighborOf`` triples. It compares TF-IDF topical edges with KG-augmented edges for synthetic query traces. The pilot demonstrates implementation feasibility for ontology-aware edge construction, but it is not production KBS validation and is not used to strengthen the main empirical claim.

Supplementary Diagnostic Figures Figure fig:redundancy-comp provides the supplementary structural diagnostic figure referenced in the main text. [pos=ht] [Figure: figures/fig_redundancy_comp.png] Redundancy density and compensation analysis. Panel (a): per-trace redundancy density distribution (mean 0.3530). Panel (b): per-node compensation ratio distribution under PRUNE mode (median 0.000, mean 0.0084). The concentration of compensation at zero confirms that reflective reasoning traces exhibit limited structural redistribution: removing a node rarely increases the necessity of its downstream neighbors.

Stratum-Dependent Held-Out Analysis The Stage 2 held-out validation results stratified by trace difficulty tier show substantial heterogeneity: `S_high` yields Spearman = 0.287 (CI [0.108, 0.451]), while `S_mid` and `S_rand` include zero in their confidence intervals. The aggregate signal (= 0.163, CI [0.092, 0.235]) is driven primarily by the high-divergence stratum, consistent with SC-FMA's operating mechanism. Full stratum results are archived in the structured data file ``outputs/stage2_stratified_metrics.json``.

Taxonomy-Stratified Ablation The ablation analysis stratified by reflective operation category reveals that structure alignment () provides the largest gain for `ERROR_CORRECTION` and `VERIFICATION` categories, while redundancy penalty () contributes most for `BACKTRACKING` and `PLANNING` categories. Full taxonomy-stratified ablation results are archived in the structured data file ``outputs/counterfactual_ablation_results.jsonl``.

Implementation Details and Reproducibility

Reproducibility Checklist The following checklist documents the measures taken to ensure full reproducibility of all experimental results reported in this paper.

Environment. Python 3.11, dependencies specified in `requirements.txt` with pinned versions. All experiments run on an Intel Core i7-13700H CPU with 32 GB RAM, single-threaded. Random seeds. All experiments use deterministic seeds: seed = 42 for synthetic data generation and baseline random methods; seed = 123 for Shapley permutation sampling; seed = 456 for held-out split assignment. NumPy, PyTorch, and Python random states are all seeded.

Synthetic data. Traces are generated via the controlled probabilistic model described in the Synthetic Data Generation section of the main manuscript. The generation script (`scripts/generate\_synthetic\_traces.py`) is deterministic given the seed. Generated traces are archived as JSON artifacts.

Hyperparameter tuning. Grid search on the validation set (40 traces) with the ranges reported in the hyperparameter tuning table of the main manuscript. No test-set information is used during tuning. The optimal configuration is frozen before test-set evaluation.

Baseline methods. All baselines use publicly available implementations: Captum for Gradient and Integrated Gradients; transformers library for attention extraction; custom NumPy implementation for MC Shapley with $M = 500$ permutations and convergence tolerance 10^{-3} . SLSQP solver. SciPy's `scipy.optimize.minimize` with method `SLSQP`, tolerance `ftol = 1e-8`, maximum iterations = 100. Warm-start from Ridge solution.

Statistical tests. Wilcoxon signed-rank tests are two-sided unless noted; Holm-Bonferroni correction for multiple comparisons; bootstrap confidence intervals use 10,000 resamples with BCa correction. Effect sizes reported as $r = Z/N$.

MuSiQue audit route. The MuSiQue route is reproduced with `python scripts/run\_knowledge\_audit\_eval.py -hf-dataset dgslibisey/MuSiQue -source-format musique -hf-split train -max-records 2000 -shuffle-seed 42 -bootstrap-samples 1000 -bootstrap-seed 42`. It writes `outputs/kbs\_real\_audit\_v1/knowledge\_audit\_report.json`, `knowledge\_audit\_summary.md`, `knowledge\_audit\_traces.jsonl`, and `dataset\_manifest.json`; the run uses zero API calls. Because the audit labels and scoring features are both deterministic functions of step type, this route is a feasibility demonstration only and is not used as positive audit-prioritization evidence.

WebQSP trace-audit diagnostic. Development uses `D: CF data raw webqsp WebQSP.train.json`; the locked diagnostic uses `D: CF data raw webqsp WebQSP.test.json`. The route is reproduced with `scripts/run\_webqsp\_trace\_audit.py`, `scripts/validate\_webqsp\_trace\_audit.py`, and `scripts/analyze\_webqsp\_separability.py`. It writes `samples.jsonl`, `reasoning\_traces.jsonl`, `replay\_results.jsonl`, `importance\_targets.jsonl`, ranking outputs, data-audit reports, and separability diagnostics under `outputs/webqsp\_trace\_audit\_v1\_train` and `outputs/webqsp\_trace\_audit\_v1\_test`. It is a supplementary diagnostic only.

Output artifacts. All raw results (CIU vectors, necessity scores, redundancy matrices, calibrated weights, ranking metrics) are saved as JSONL or JSON artifacts in the `outputs/` directory. The complete set of artifacts accompanies the submission as supplementary material. The MuSiQue route is archived under `outputs/kbs\_real\_audit\_v1` and is reported as a constructed-label feasibility demonstration (`kbs\_style\_audit\_prioritization\_evidence\_only`); the WebQSP route is archived under `outputs/webqsp\_trace\_audit\_v1\_test` and is reported as a fixed-schema separability diagnostic.

SLSQP Convergence Diagnostics SLSQP convergence statistics across $N = 1,027$ steps (200 traces): mean iterations to convergence 5.3 2.1 (std), median 5, maximum 12. The worst-case trace ($k=8$, dense redundancy) required 12 iterations and 14.7 ms. Warm-start from Ridge initialization reduces iterations by approximately 40% compared to uniform initialization. Under default hyperparameters ($=1.0$, $=0.5$, $=0.2$, $=0.1$, $=10^{-4}$), the Hessian condition number (H) 3.2, guaranteeing numerically stable optimization. Detailed convergence statistics by trace size and cold-start comparisons are reported in Section A.6 below.

Graph Construction Parameter Sensitivity The keyword-heuristic graph construction introduces two free parameters: $_tfidf = 0.35$ and $w_topical = 0.75$. Sensitivity analysis (20% perturbation) shows $|| < 0.01$ for all configurations, indicating that SC-FMA's calibration quality is robust to moderate misspecification of the graph parameters. Full sensitivity tables and per-trace-size breakdowns are archived in the structured data file ``outputs/phase6_sensitivity.json``.

Controlled Proxy-Label Independence Analysis We verify that the controlled proxy-label definition used in the synthetic benchmark does not share input features with the SCU objective. The correlation between each proxy-label component and each SCU input is summarized in Table `tab:oracle-independence`.
Cross-correlation between controlled proxy-label components and SCU inputs across 1, | c_i (CIU) | n_i (necessity) | r_i (redundancy) s_i (correctness) | 0.08 | -0.04 | -0.07 d_i (lexical diversity) | -0.12 | 0.06 | 0.03 p_i (position value) | 0.11 | -0.09 | -0.02 [TABLE]

All cross-correlations are below	< 0.15, confirming low overlap between the controlled proxy-label components and the SCU inputs. The previous isomorphic definition ($y_i = 0.4c_i + 0.4n_i + 0.2(1-r_i)$) achieved P_e
----------------------------------	---

Verification Boundary

MuSiQue KBS-style Knowledge-Audit Details

`kbs_style_audit_prioritization_evidence_only`

Hyperparameter sensitivity is reported as a supplementary QP robustness diagnostic only; it is not positive external validation and does not change the PRM800K Ridge recommendation.